

# Masked Self-Attention

## Masking the future in self-attention

- To use self-attention in **decoders**, we need to ensure we can't peek at the future.
- At every timestep, we could change the set of **keys** and **queries** to include only past words. (Inefficient!)
- To enable parallelization, we **mask out attention** to future words by setting attention scores to  $-\infty$ .

$$e_{ij} = \begin{cases} q_i^\top k_j, j \leq i \\ -\infty, j > i \end{cases}$$

For encoding these words

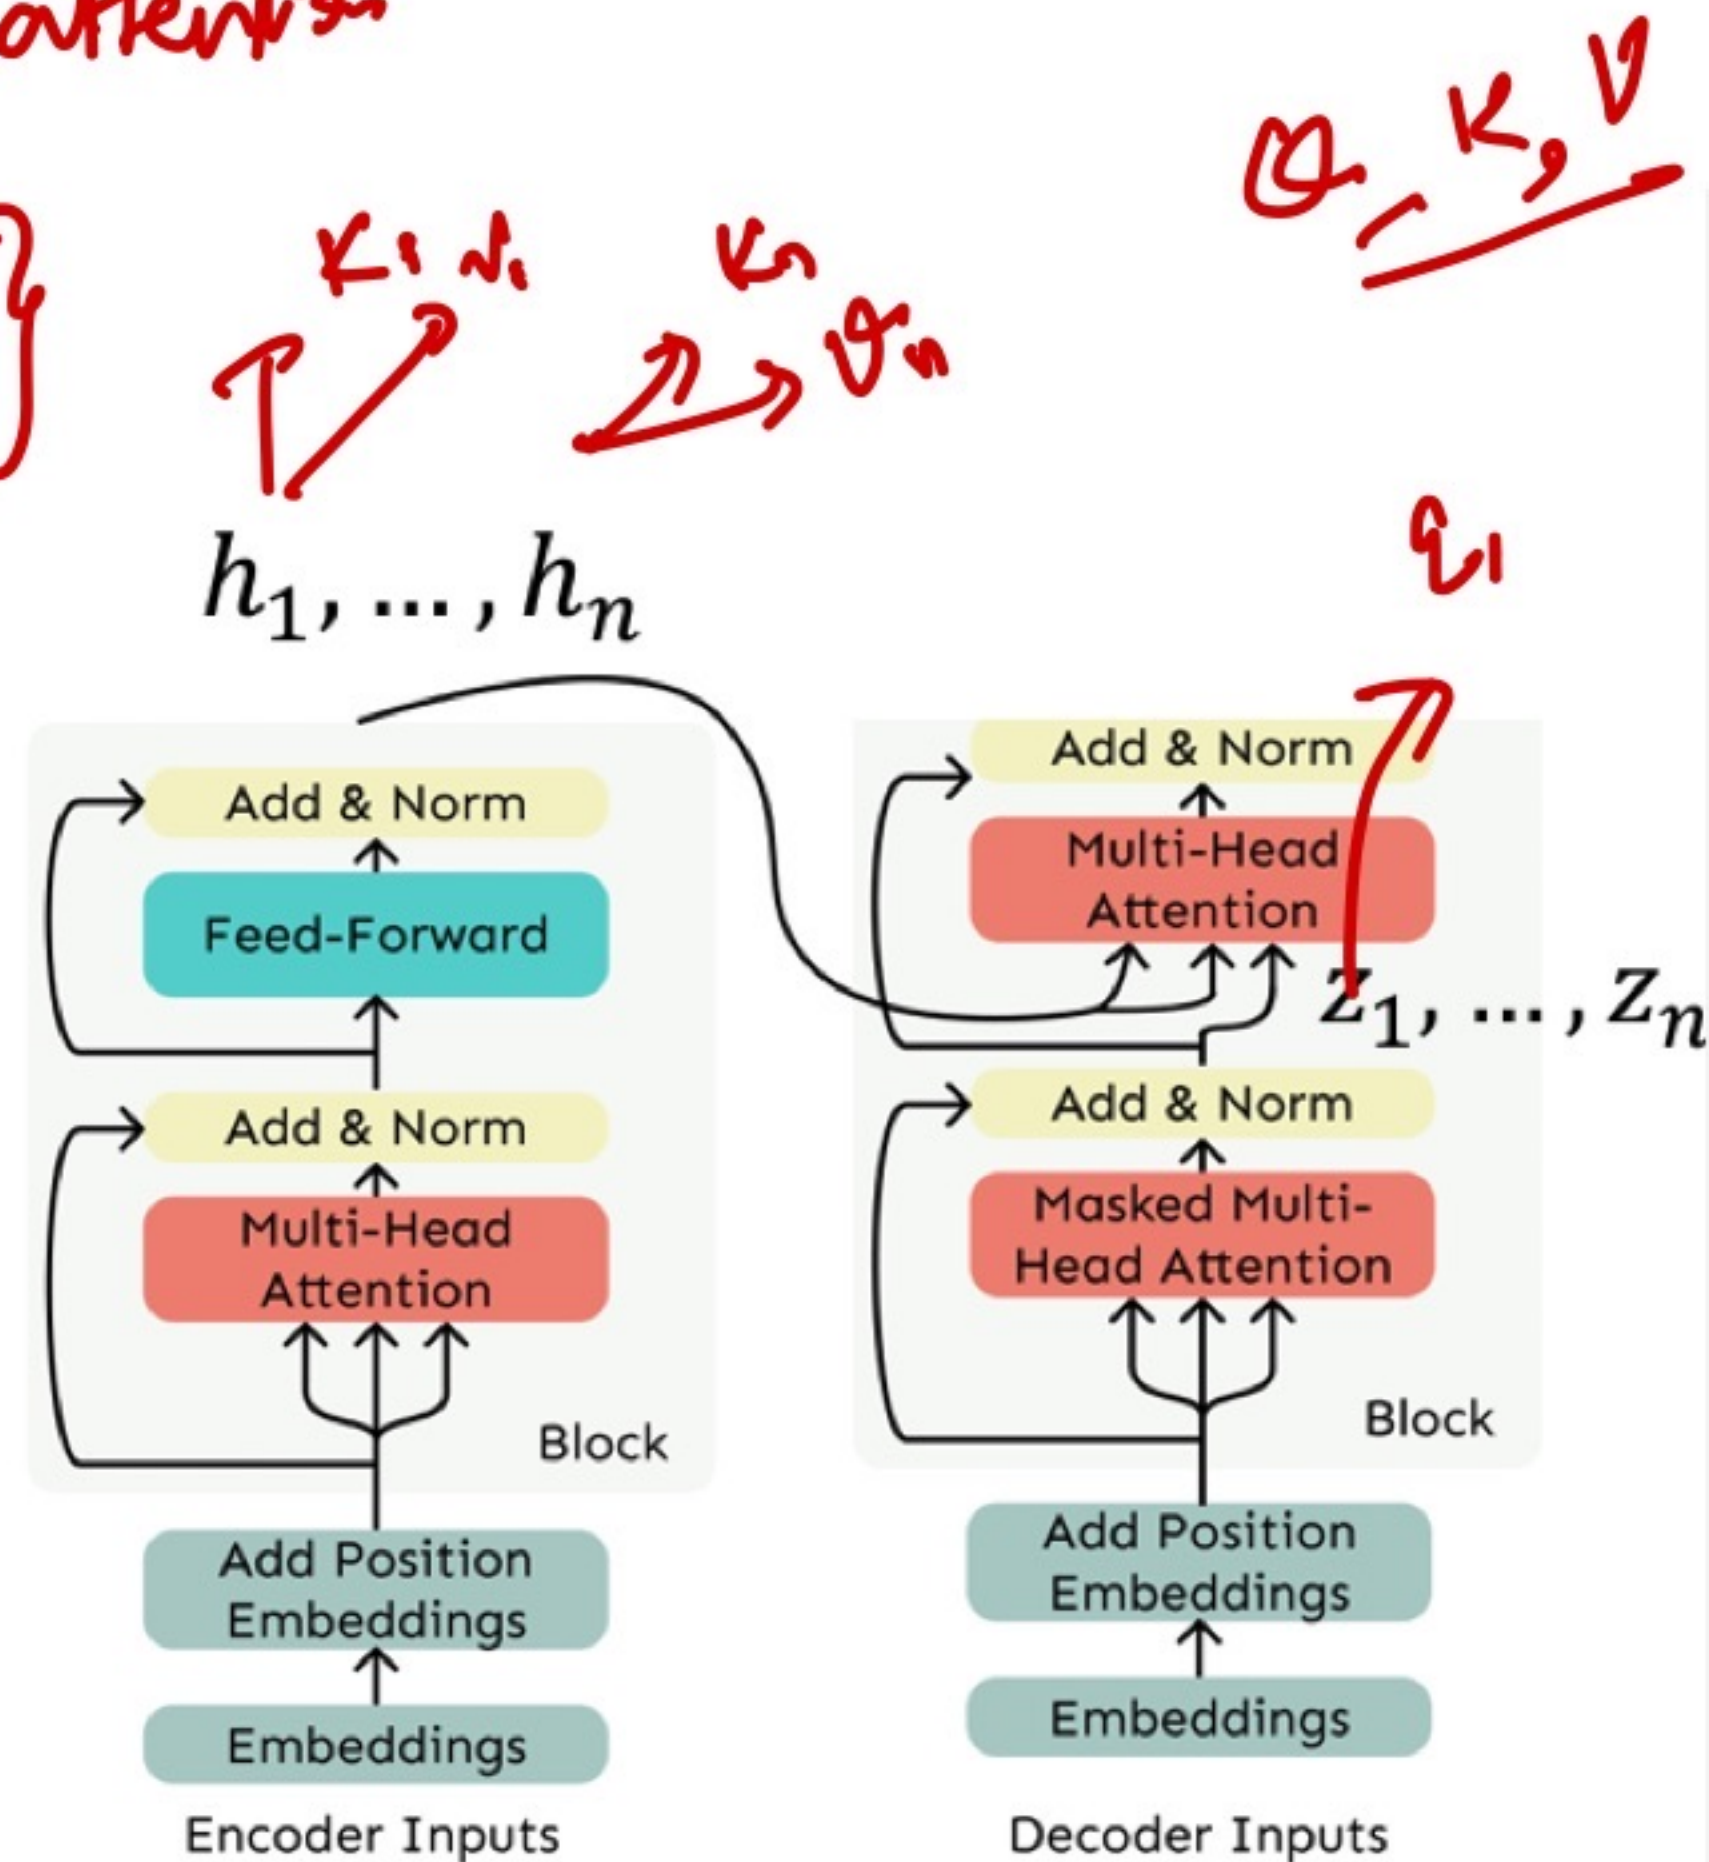




# Encoder-Decoder Attention

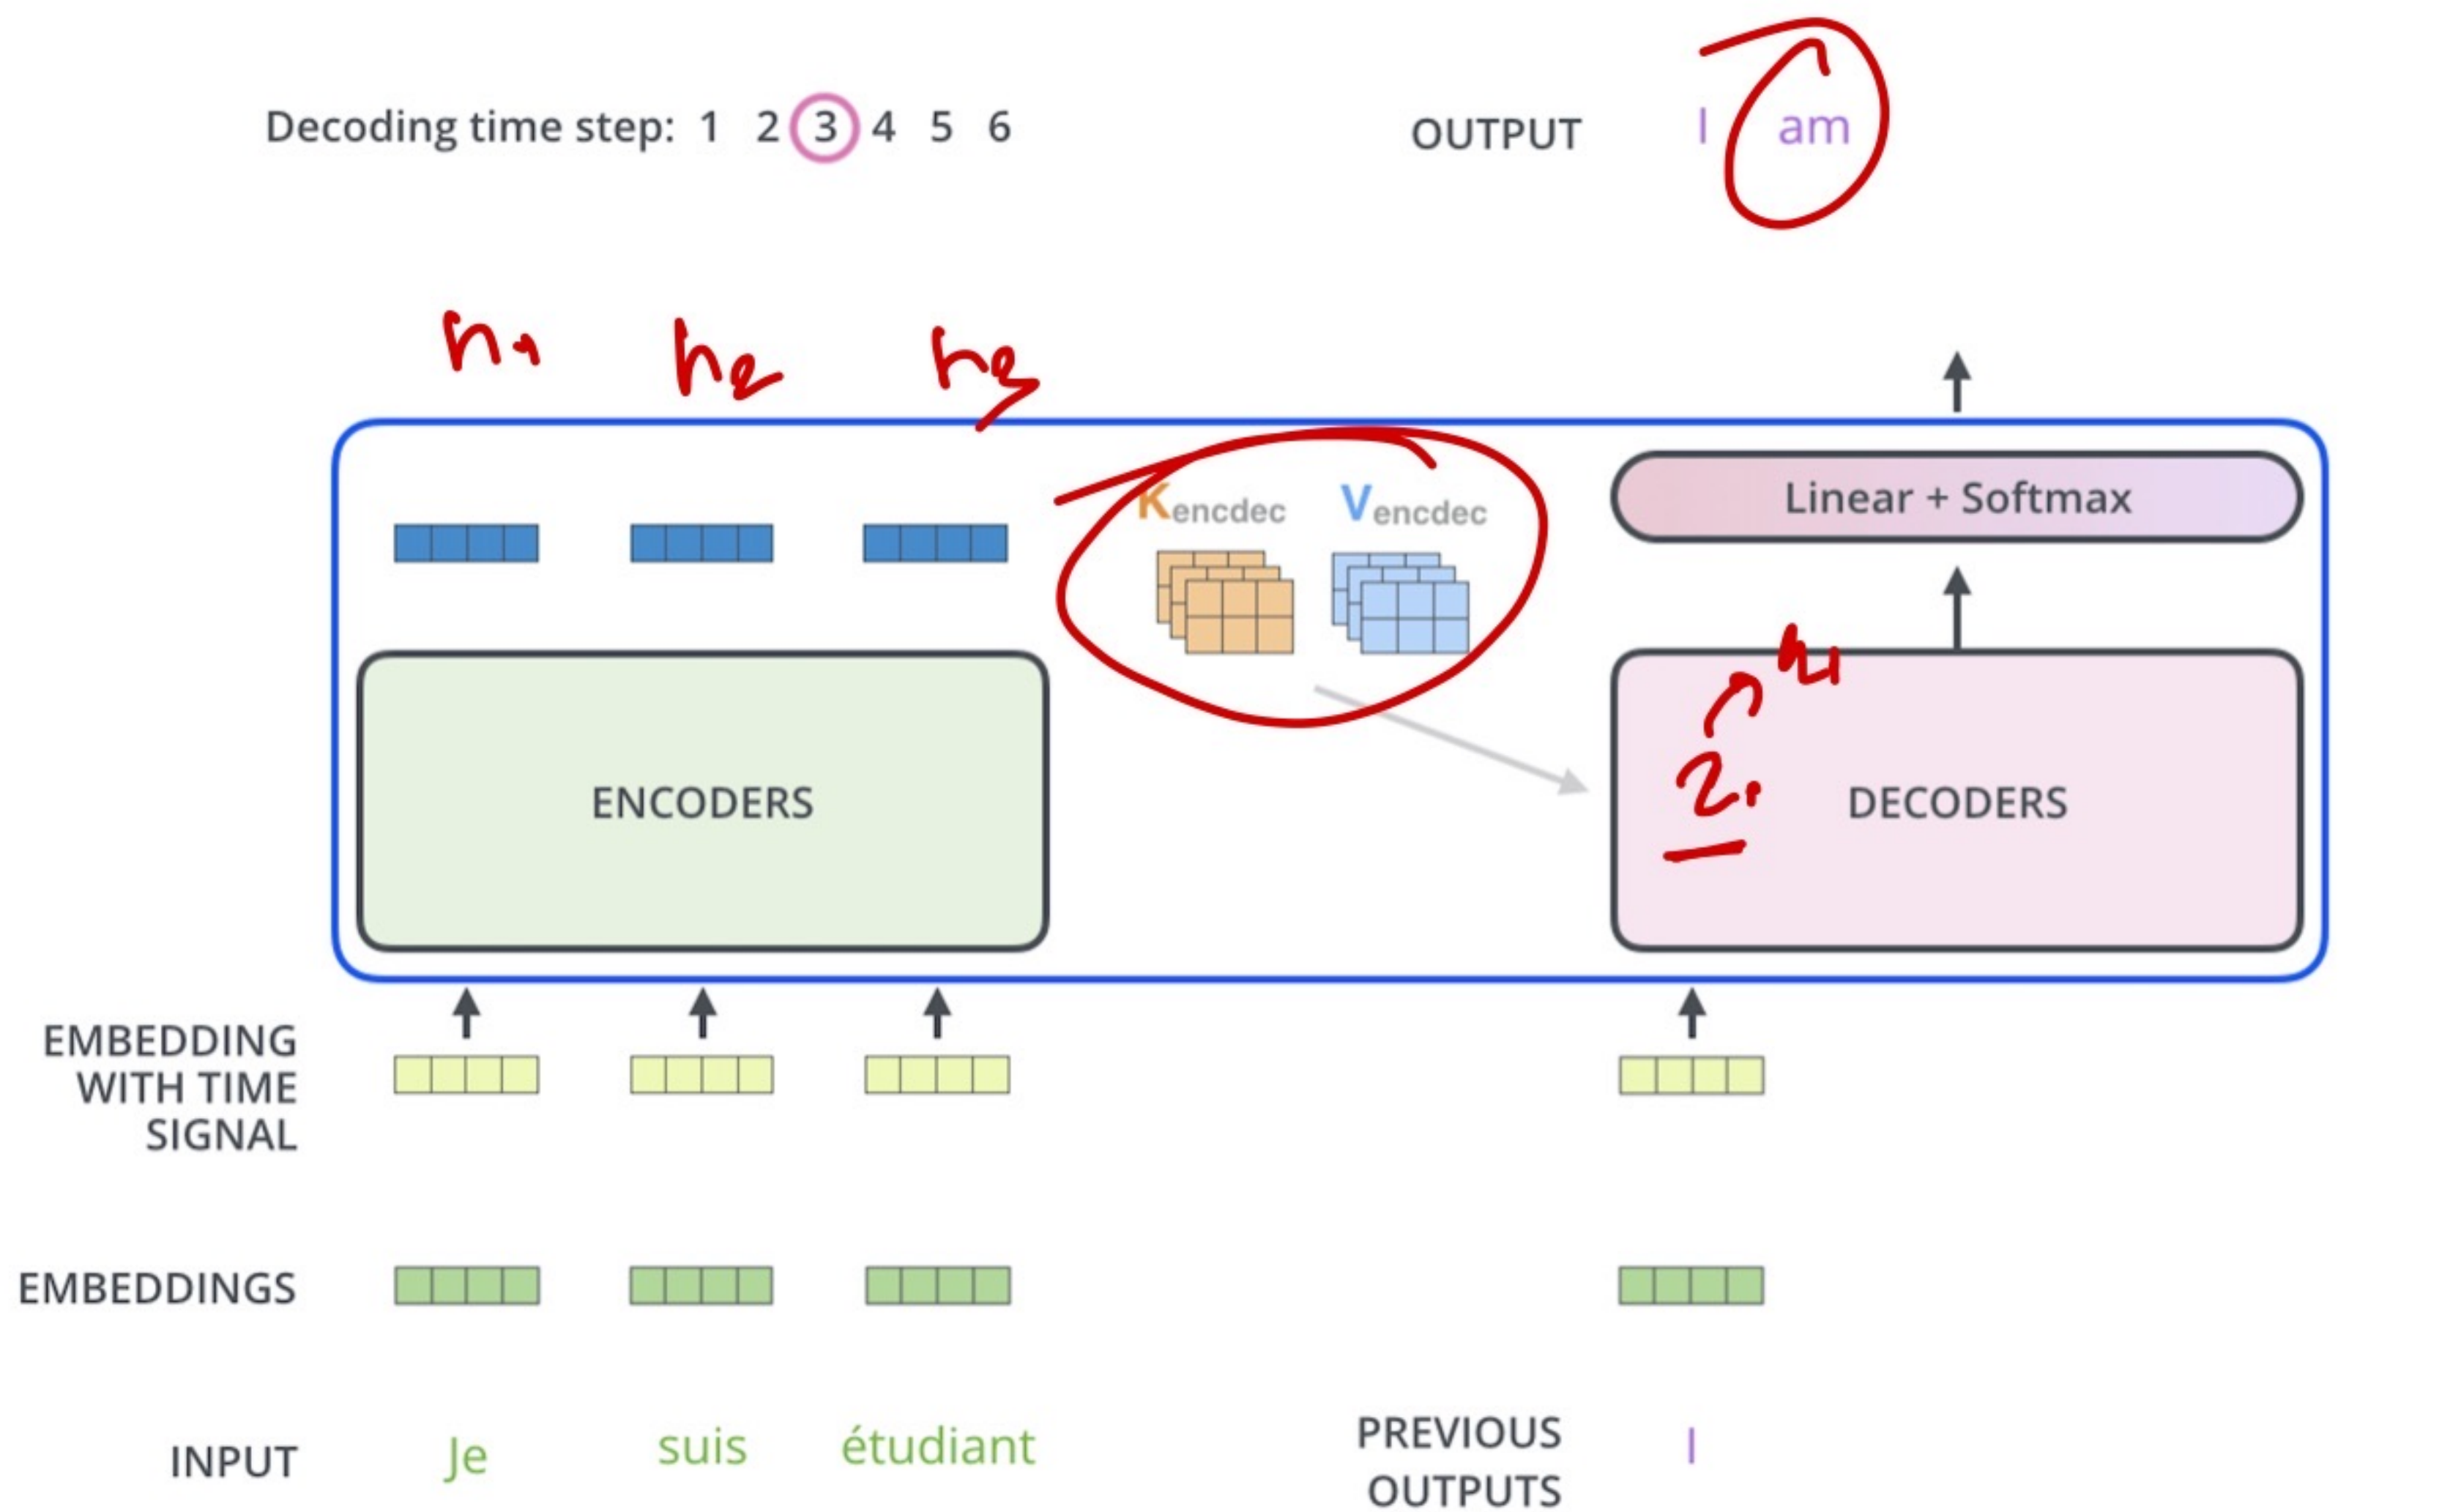
*cross-attention*

- We saw that self-attention is when keys, queries, and values come from the same source.
- In the decoder, we have attention that looks more like what we saw last week.
- Let  $h_1, \dots, h_n$  be **output** vectors from the Transformer **encoder**;  $x_i \in \mathbb{R}^d$
- Let  $z_1, \dots, z_n$  be input vectors from the Transformer **decoder**,  $z_i \in \mathbb{R}^d$
- Then keys and values are drawn from the **encoder** (like a memory):
  - $k_i = Kh_i, v_i = Vh_i$ .
- And the queries are drawn from the **decoder**,  $q_i = Qz_i$ .

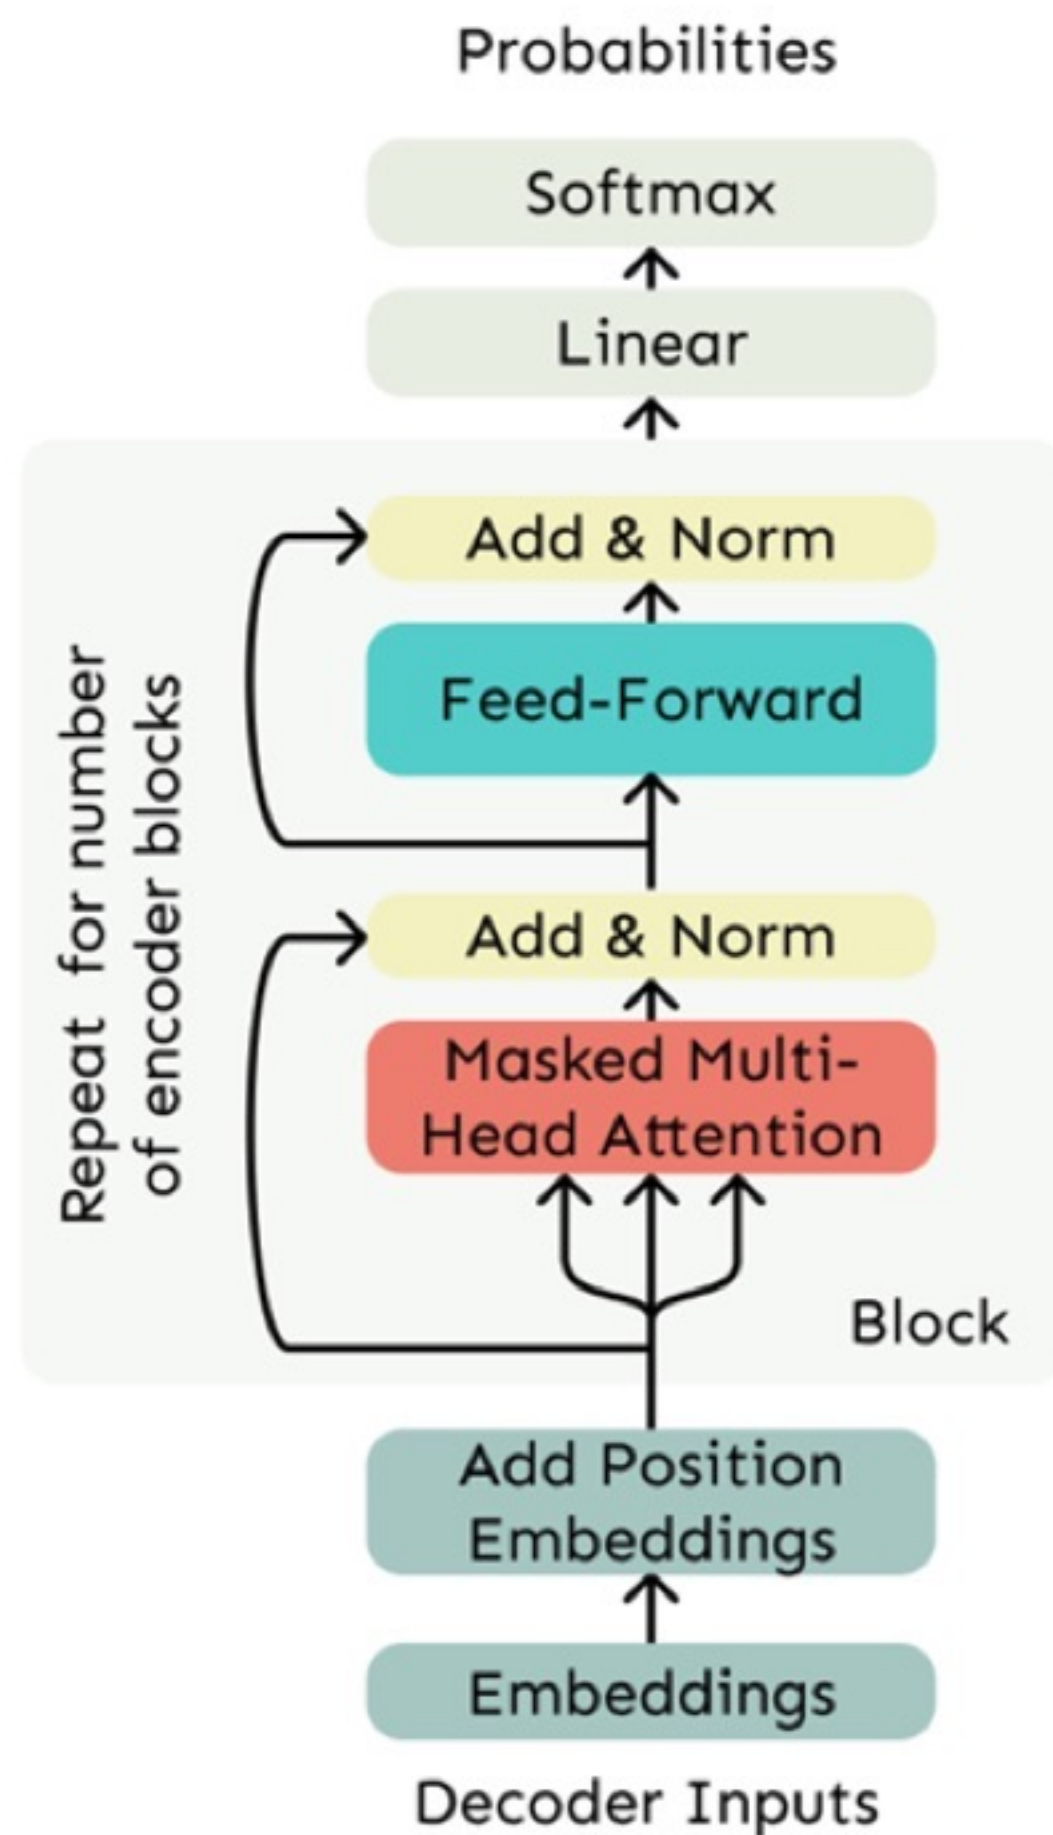




# Encoder-Decoder Attention

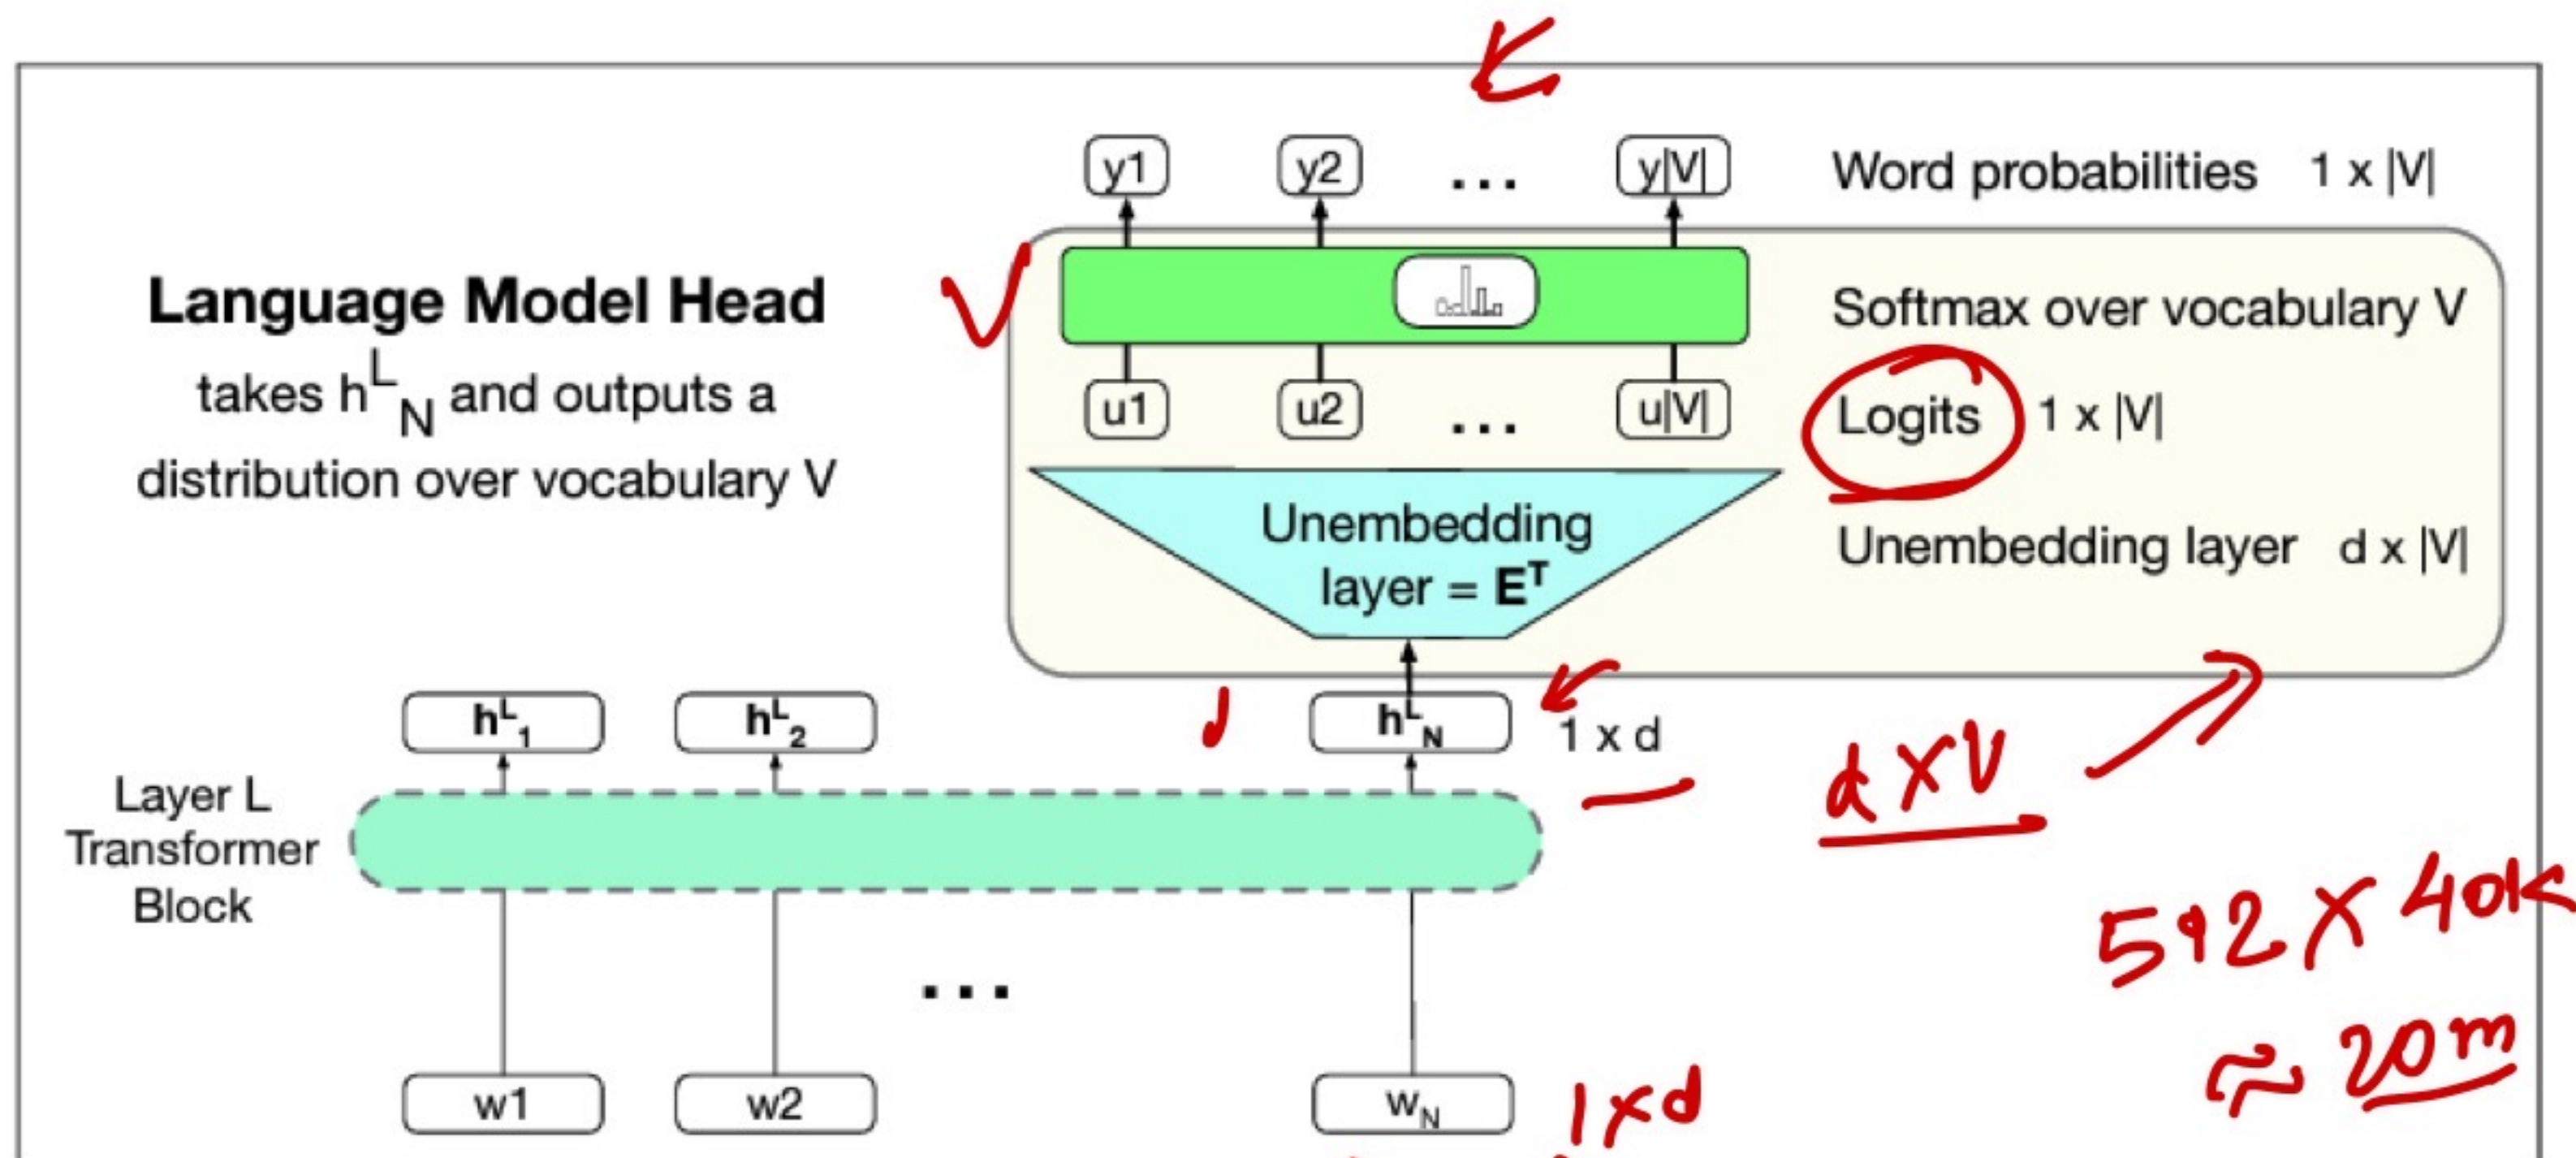


# Transformer as a Language Model (Decoder)





# Unembedding layer in Transformer LM



**Figure 10.13** The language modeling head: the circuit at the top of a transformer that maps from the output embedding for token  $N$  from the last transformer layer ( $h_N^L$ ) to a probability distribution over words in the vocabulary  $V$ .

# Weight Tying


- The unembedding layer can be learned, but it is also very common to tie this matrix to the embedding matrix  $E$ .
- At the input, embedding layer ( $[E : V \times d]$ ) is used to map from one-hot ( $V$ -dim) to an embedding ( $d$ -dim).
- At the language modeling head, transpose of the embedding layer ( $[E^T : d \times V]$ ) is used to map back from the embedding ( $d$ -dim) to a vector over the vocabulary ( $V$ -dim).
- In the learning process,  $E$  is optimized to be good at doing both of these mappings.



# Text Completion via Language Models

## Conditional Generation Task

The task of generating text conditioned on an input piece of text



**Figure 10.15** Autoregressive text completion with transformer-based large language models.

# Large Language Models: Main Insight

- LM is given a test ~~suffix~~ <sup>prefix</sup> (context) and is asked to generate a possible completion
- As the generation proceeds, model has access to the context as well all of its previously generated tokens
- This ability is the key to the power of Large Language Models built from transformers

*But why do we care about predicting next words?*

Many practical NLP tasks can be cast as word prediction / text completion

